



Other Balanced Trees, and an Unbalanced One (Huffman Encoding)

EECS 233

Previous Lecture

- AVL trees (self-balancing binary tree)
- Rotation operations: left rotation, right rotation, double rotation



Previous Lecture

- Insertion may cause change of balance value
- Insertion algorithm
 - Insert a node, which may cause a change to balance values
 - For each ancestor (from the leaf to the root), if the balance value
 - ✓ changes from ± 1 to 0, DONE
 - ✓ changes from 0 to ± 1 , go to the next ancestor
 - ✓ changes from ± 1 to ± 2 , rebalance this tree, DONE
 - How to balance the tree? If insertion occurred in the
 - ✓ left-left subtree, do right rotation
 - ✓ right-right subtree, do left rotation
 - ✓ left-right subtree, do left-right double rotation
 - ✓ right-left subtree, do right-left double rotation

Deletion from AVL Trees


- Deletion may also cause unbalanced subtrees; if a deletion causes a change to the height, the ancestors need to be checked or balanced.
- Deletion algorithm
 - remove the node (leaf or non-leaf, **remember how it works?**)
 - Ultimately causes removal of a node X with 0 or 1 child
 - May cause a change of the balance values for X's parent
 - For each ancestor (from X's parent to the root), if the balance value
 - ✓ changes from ± 1 to 0, **what to do?**
 - ✓ changes from 0 to ± 1 , **what to do?**
 - ✓ changes from ± 1 to ± 2 , **what to do?**
 - To perform the balancing, use the same rotations as in insertion

Efficiency of AVL Trees

- An AVL tree containing n items has a height that is $O(\log_2 n)$.
- Search and insertion are both $O(\log_2 n)$.
 - Search travels at most one path down the tree
 - An insertion goes down one path to the insertion point, and then goes back up adjusting balances/performing rotations
 - ✓ in the worst case, both the path down and the path back up consider $O(\log_2 n)$ nodes
 - ✓ each rotation requires a constant number of operations, and we perform at most two of them for a given insertion
- Deletion begins with the same procedure used in binary search trees ($O(\log_2 n)$ complexity). It can also require rotations to rebalance the tree, also at the complexity of $O(\log_2 n)$.

Other Trees

- 2-3 trees
- Splay trees
- Red-black trees
- B-trees




Aren't AVL Trees Perfect?

■ AVL tree - binary balanced tree

- Searching an item requires traversing from root to (at most) leaf
- $\log N$ descents (random accesses)
- Visiting each node - one comparison.

■ On disk:

- Random access is expensive
- Disk read is in blocks (e.g. 4K bytes at once)
- Disk-based AVL trees waste much of a block for an internal node read
 - ✓ The node is smaller than the block
 - ✓ Payload of the node is not used



B-tree motivations

- not all operations take equal time
 - CPU: 3 billion operations per second
 - Hard Disk access: ~9 ms to read a block
⇒ ~111 accesses per second
 - (newer SSDs: 0 ms seek time, latency .09 ms
⇒ have to design new search trees to match new characteristics)
- Each disk access = 27 million CPU operations

Search trees in the real world

■ Typical large database:

- 10 million records
- each key: 32 bytes (e.g. name)
- each record: 256 bytes
- ~ 3 TB of data

■ How much time to we have to access?

- We're one of many users on system, say 30
- in 1 sec we get 100 million CPU operations and about 4 disk accesses.

What can we do with what we have?

- 100 million operations and 4 disk accesses
- unbalanced binary tree:
 - 10 million disk accesses * 4/sec = a *lot* of time (1.24 years!)
- perfectly balanced binary tree
 - On average, need $1.38 \log N$ disk accesses
 - $\log_2 10 \text{ million} \approx 24$
 - So about 32 disk access total = 8 seconds
- typical binary tree: a few nodes are several times deeper
 - Could require several times 8 seconds
- AVL tree: typical case very close to $\log N$
 - Still 6 seconds for each search



What can we do?

- Design better search tree:
 - requires no more than 3 or 4 disk access
 - don't care (much) about CPU operations
 - essentially free compared to disk accesses

- can make search and data structures more complex provided we minimize the disk accesses

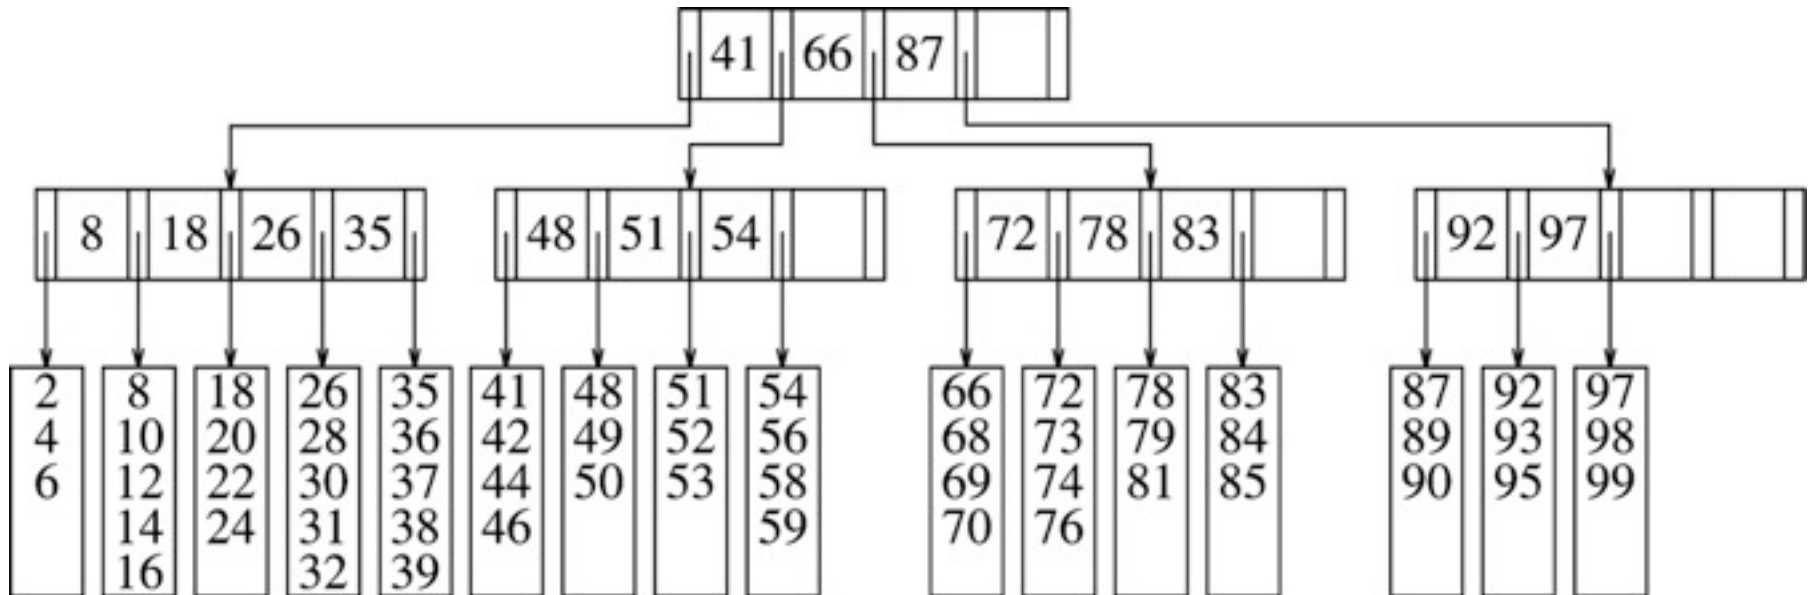
- Conclusion:
 - can't have deep trees
 - need wide, shallow trees

Designing for Disk Particulars



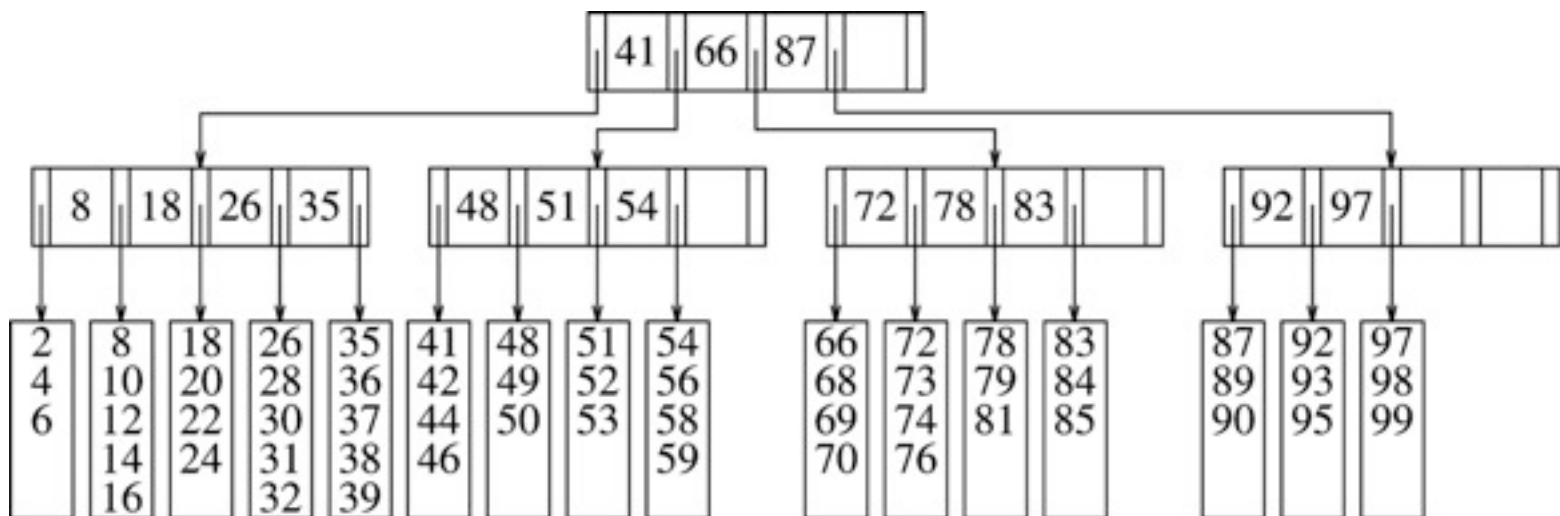
B-Tree Example

- Order: branching $M=5$, items in leaf $L=5$
- Example: search 74



B-Tree Example

- A B-tree of order M is a balanced tree in which
 - Each non-leaf node stores up to $M-1$ keys and M pointers; key i represents the smallest key in subtree $i+1$ (pointer $i+1$)
 - Data items are stored in leaf nodes (L =maximum number of items.)
 - Each node except for root is at least half full
 - ✓ Non-leaf node: at least $M/2$ pointers
 - ✓ Leaf node: at least $L/2$ entries
 - The root is either a leaf, or has between 2 and M children
- Determining M and L
 - $(M-1) * \text{sizeof}(\text{key}) + M * \text{sizeof}(\text{pointer}) \leq \text{Block size (e.g. 4KB)}$
 - $L * \text{sizeof}(\text{item}) \leq \text{Block size}$



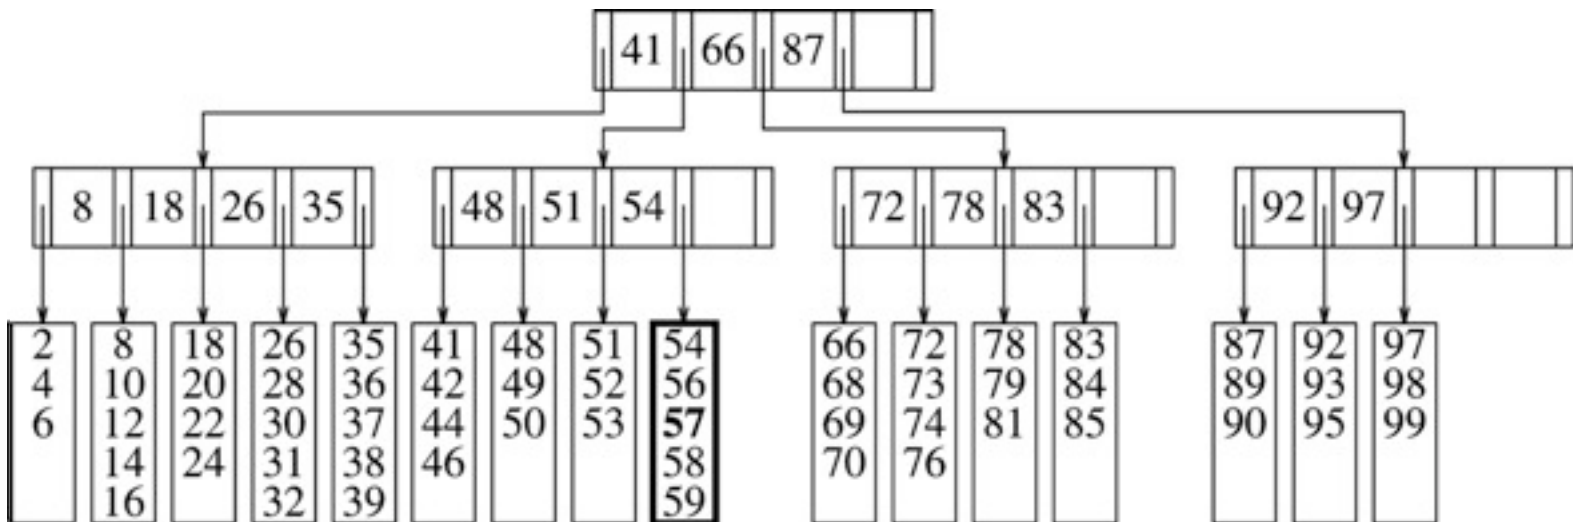
Determine B-tree parameters

- Example: 10 million records; 8k = 8,192 byte blocks
 - block size depends on application (e.g. names vs photos)
- Each node is a block
- B-tree of order M has
 - M-1 keys: each key is 32 bytes (e.g. a name) = $32 \cdot M - 32$ bytes
 - M branches: each branch is another block (4 bytes) = $4 \cdot M$ bytes
- How large can M be?
 - $36 \cdot M - 32 < 8,192 \Rightarrow M = 228$
- Suppose each record is 256 bytes
 - can fit $8192/256 = 32$ records in one block
- B-tree parameters
 - each branch node branches at least $M/2 = 114$ ways
 - between 16 and 32 records per block at leaf nodes

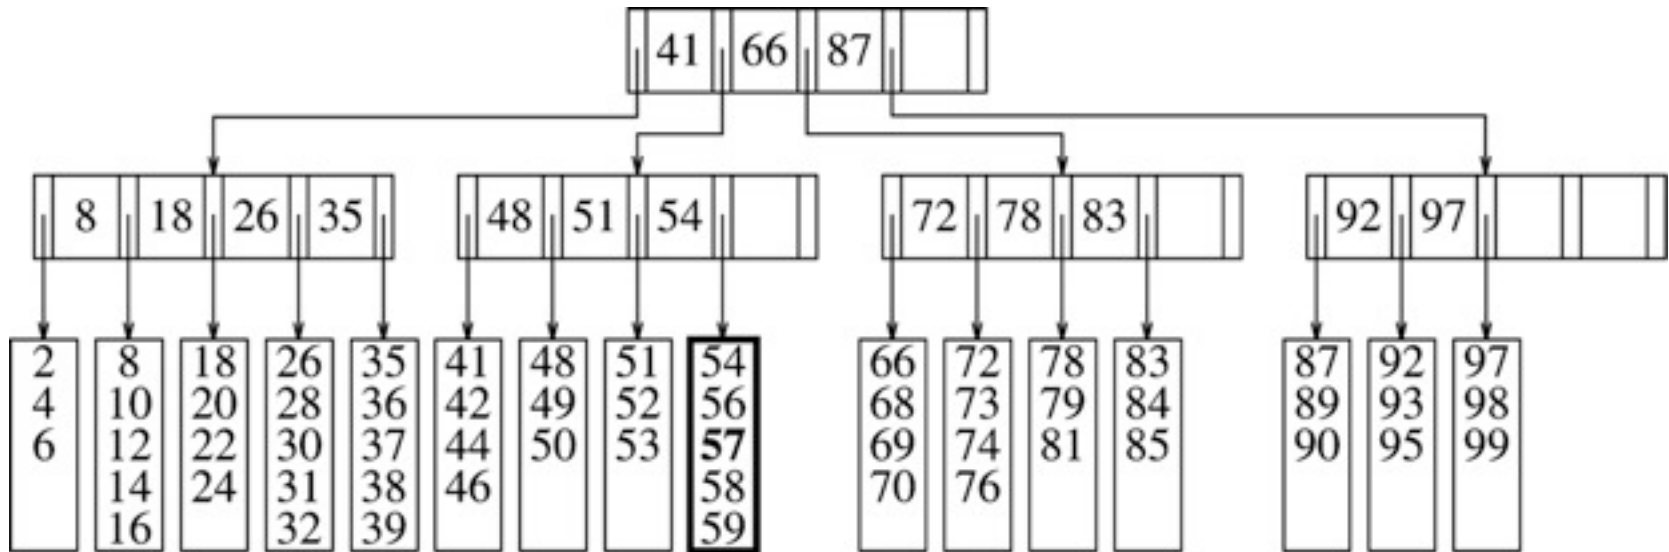
Insertion into B-Trees

- First search for the key until we reach a leaf
- Insert the key into the array of the leaf node
- ... assuming there is room

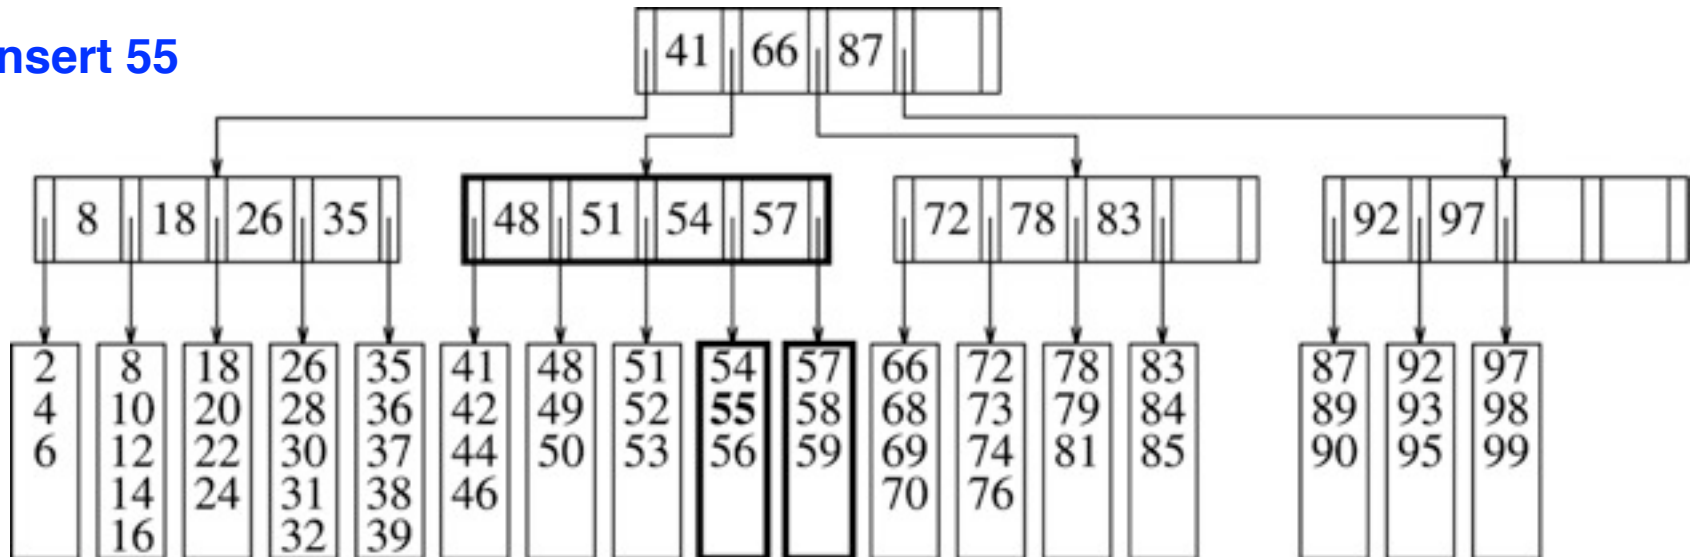
Insert 57



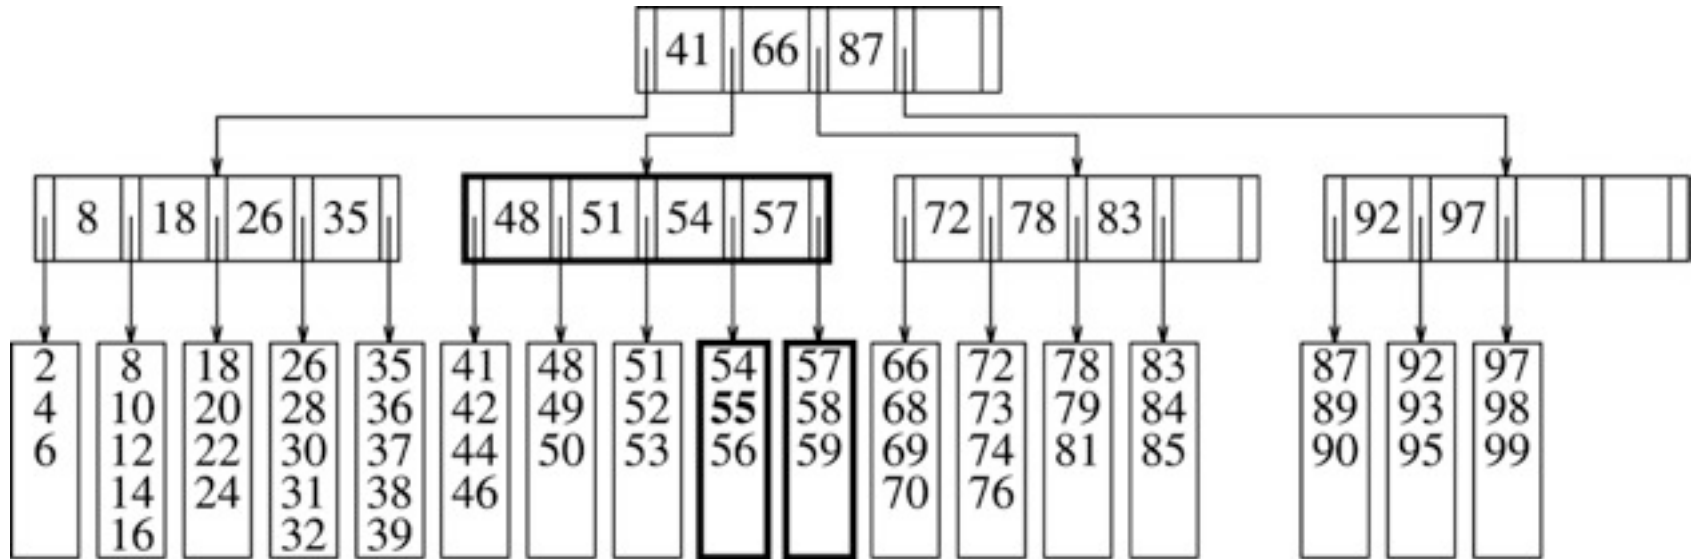
Leaf Split due to Insertion



Insert 55

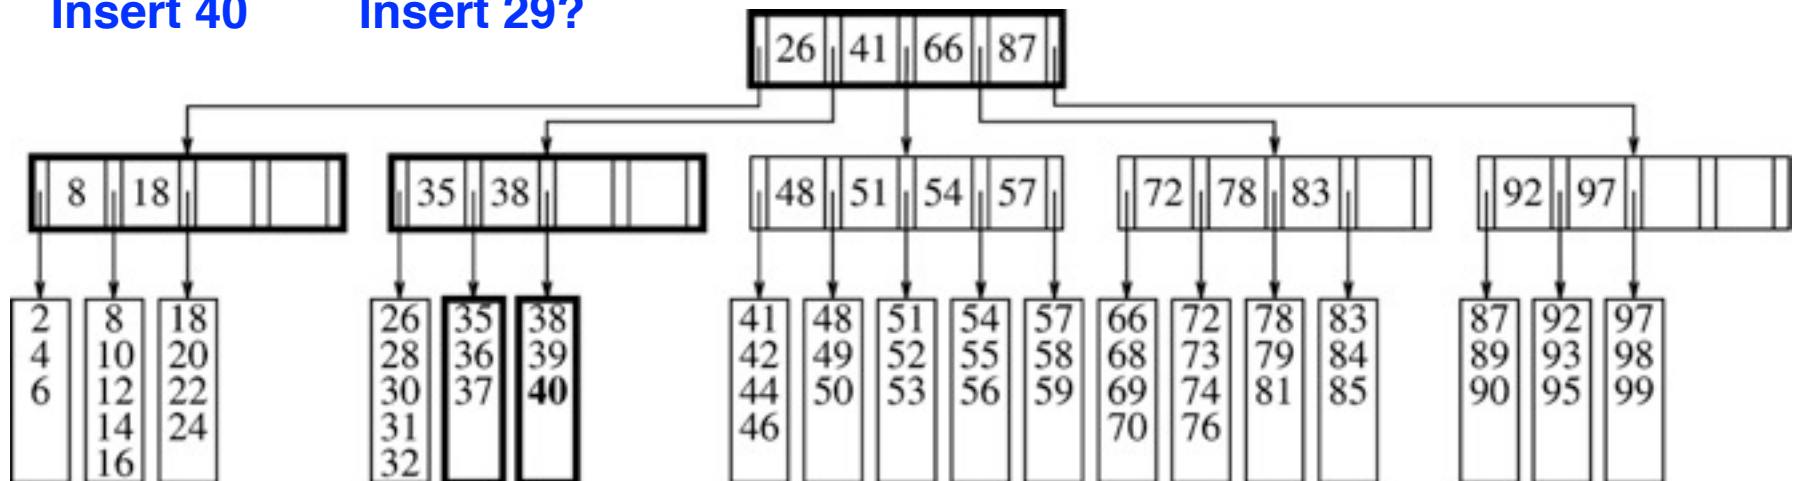


Splitting an Internal Node



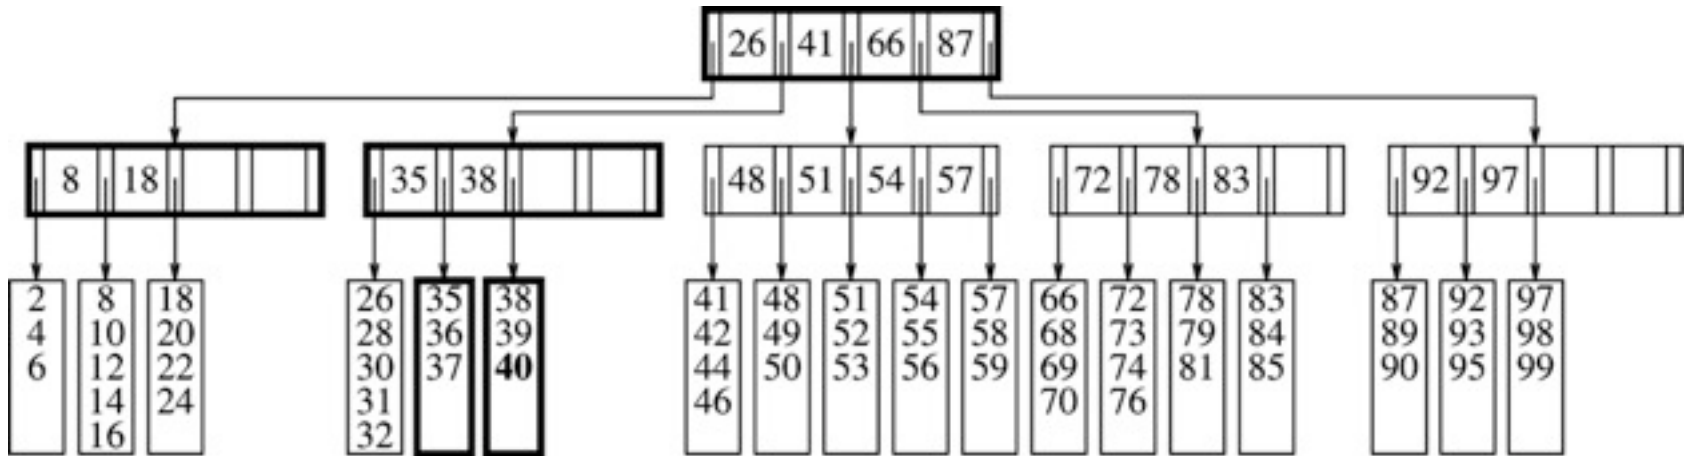
Insert 40

Insert 29?

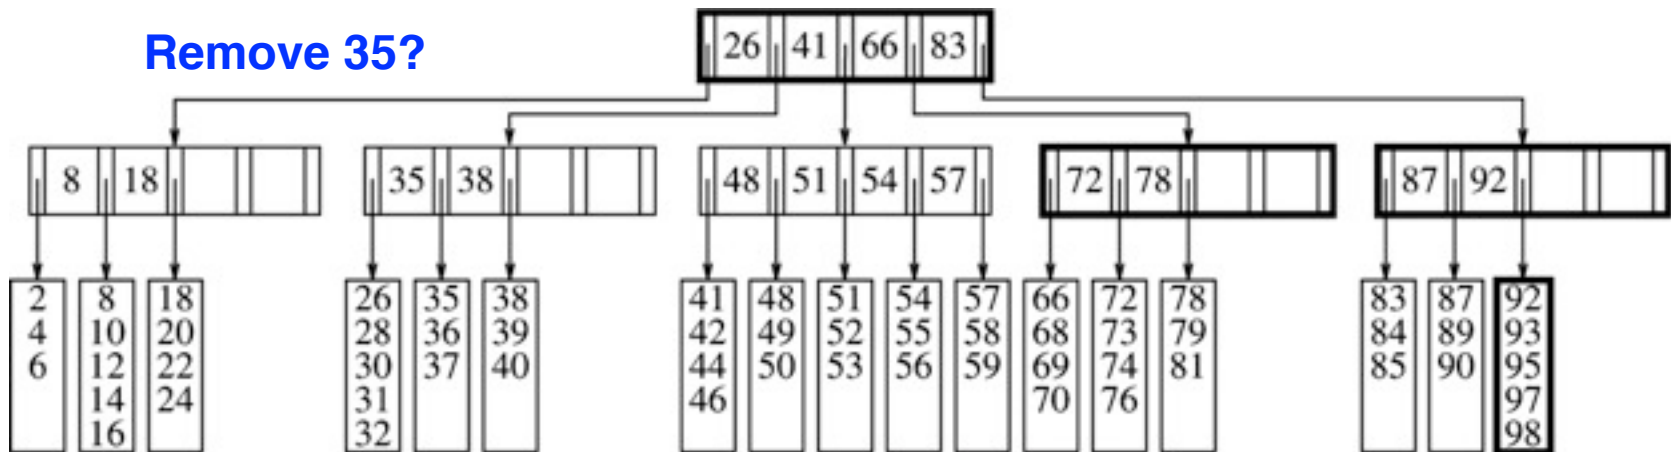


Deletion from B-trees: Node Merging

Remove 99:




Remove 35?



Efficiency of B-Trees

- Often measured by number of disk accesses, instead of CPU comparisons
- How many disk accesses are required for each search/insertion/deletion?
- CPU running time. Assume we use sorted array as the internal data structure of the nodes
 - for search?
 - for insertion/deletion?



Balanced or Unbalanced?

- Balanced search trees provide bounded time complexity
 - AVL trees
 - B-trees

Many applications require quick access to the tree data structures, and balanced trees offer guaranteed performance

- In other applications, we may not care if the tree is balanced or not, as our objective is not to save CPU time
- Compression is an example (save disk storage space, save the size of network transfer, etc)
- Example: Huffman encoding

Fixed-length Character Encoding

- A character encoding maps each character to a number.
- Computers usually use fixed-length character encodings.
 - ASCII (American Standard Code for Information Interchange) uses 8 bits per character.
 - A : 01000001
 - B : 01000010
 - C : 01000011
 - ...
 - We can represent this encoding with a binary tree (height=8)
 - Unicode uses 16 bits per character to accommodate foreign-language characters. (ASCII codes are a subset.)
- Fixed-length encodings are simple, because
 - all character encodings have the same length
 - a given character always has the same encoding

Variable-length Character Encoding

- Problem: fixed-length encodings waste space.
- Solution: use a variable-length encoding.
 - use encodings of different lengths for different characters
 - assign shorter encodings to frequently occurring characters (the characters may have very different frequencies; this is a good property that we can capitalize on)
 - E : 01
 - O : 100
 - S : 111
 - T : 00
- TEST would be encoded as 000111100
- Challenge: when decoding/decompressing an encoded document, how do we determine the boundaries between characters?
 - example: for the above encoding, how do we know whether the next character is 2 bits or 3 bits?
- Solution: no character's encoding can be the prefix of another character's encoding (e.g., couldn't have T: 00 and Q: 001).


Huffman Encoding

- Huffman encoding is a type of variable-length encoding that is based on the character frequencies in a given document.
- Huffman encoding uses a binary tree:
 - to determine the encoding of each character
 - to decode an encoded file – i.e., to decompress a compressed file, putting it back into ASCII
- Example of a Huffman tree (for a text with only six chars):

Leaf nodes are characters.


Left branches are labeled with a 0,
and right branches are labeled with a 1.

If we follow a path from root to leaf,
we get the encoding of the character in
the leaf. An example: 101 = 'i'




Building A Huffman Tree

- Begin by reading through the text to determine the frequencies.
- Create a list of nodes that contain (character, frequency) pairs for each character that appears in the text.




Remove and “merge” the nodes with the two lowest frequencies, forming a new node that is their parent.

- left child = lowest frequency node
- right child = the next lowest frequency node
- frequency of parent = sum of the frequencies of its children




Building A Huffman Tree

- Merge the two remaining nodes with the lowest frequencies:
 - This is a **greedy** algorithms; a greedy algorithm attempts to get the optimal solution by finding currently the best solution in each step



Building A Huffman Tree

- It continues



Building A Huffman Tree

- The final tree, and character encoding

